




Continued

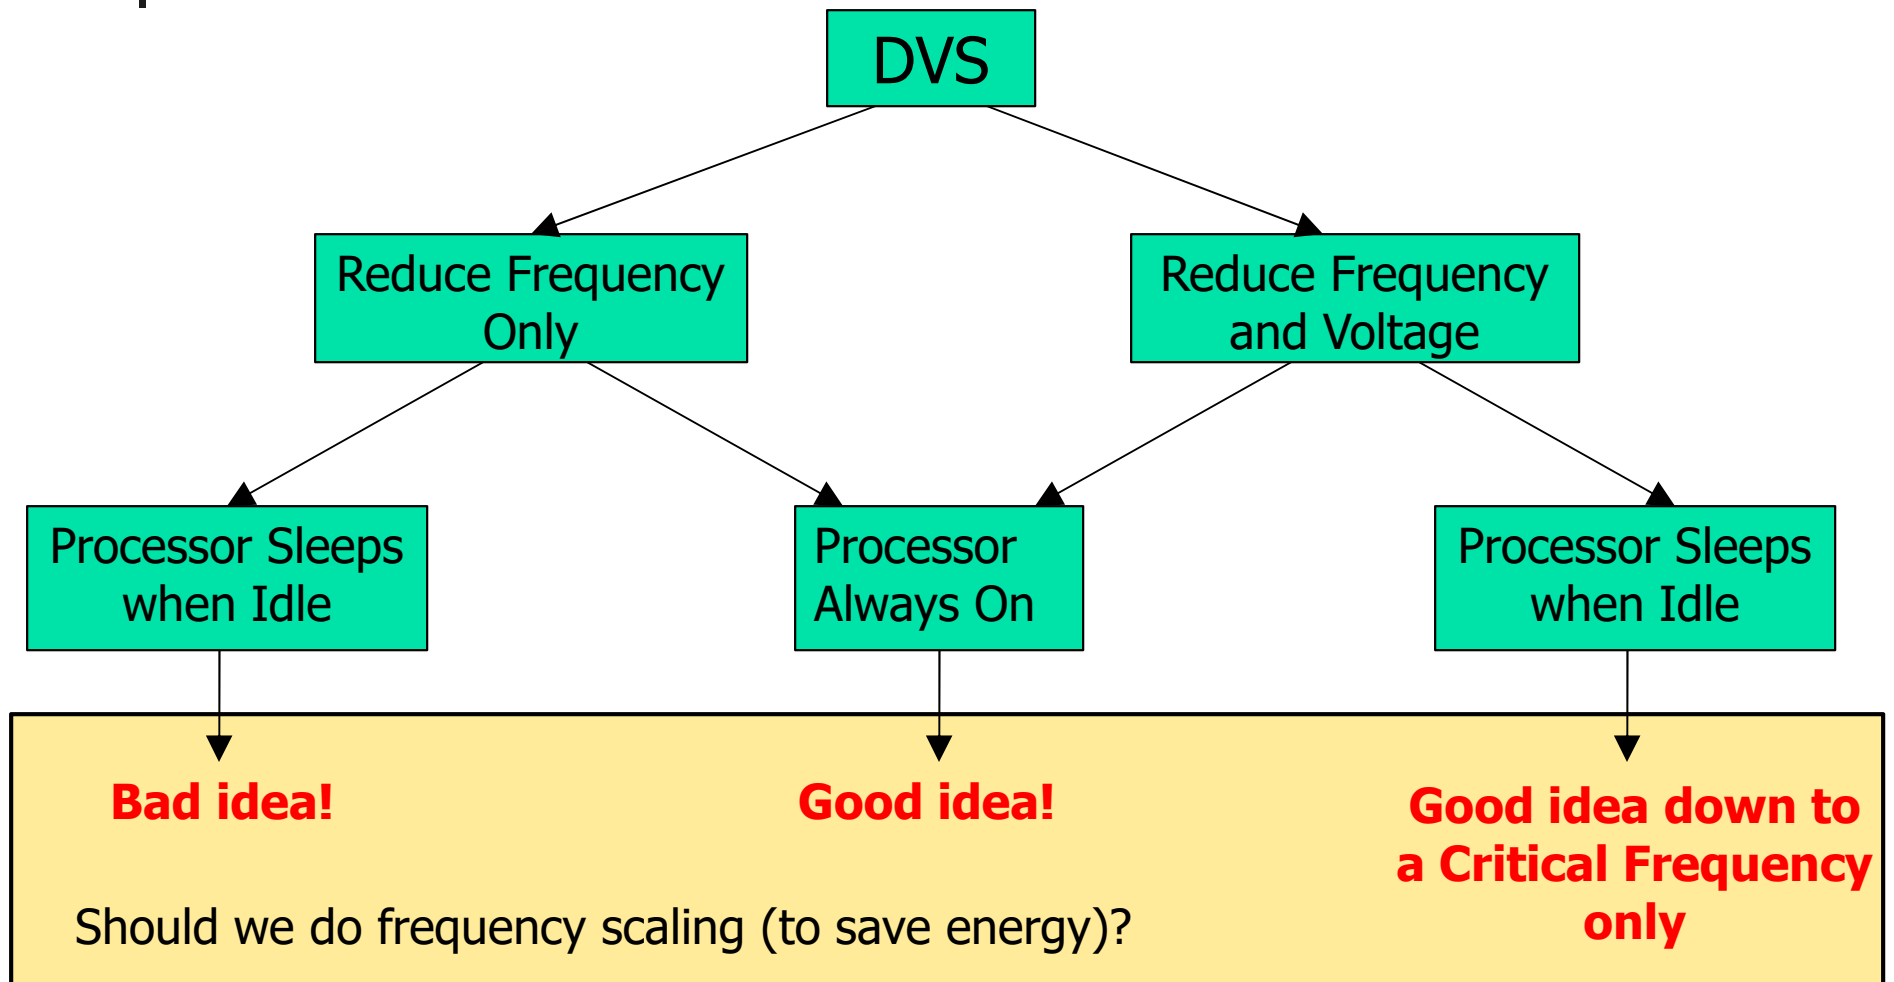
Midterm

- Average grade = 18

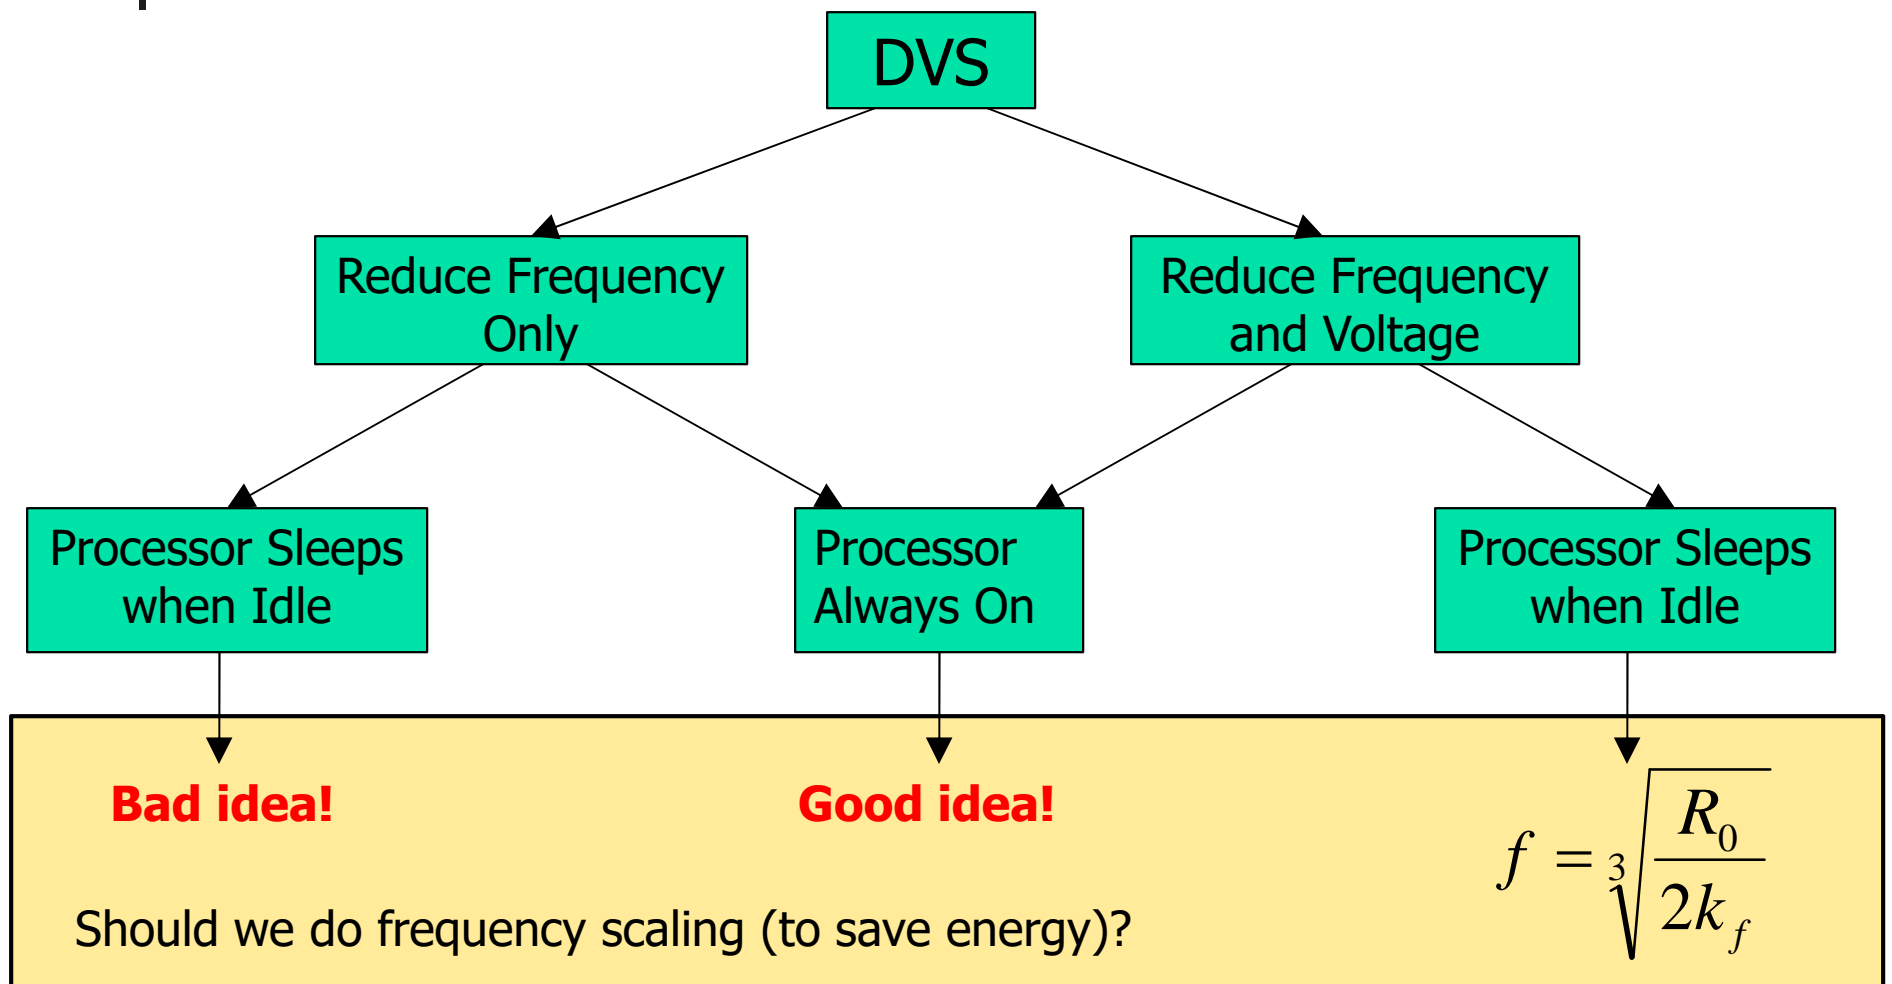




Recap



Recap





Advanced Configuration and Power Interface (ACPI)

- Defines different power saving states in a platform-independent manner
- The standard was originally developed by Intel, Microsoft, and Toshiba (in 1996), then later joined by HP, and Phoenix.
- The latest version is "Revision 6.1," published by UEFI (March 2016).



Global States

- **G0:** *working*
- **G1:** *Sleeping and hibernation* (several degrees available)
- **G2:**, *Soft Off*. almost the same as G3 *Mechanical Off*, except that the power supply still supplies power, at a minimum, to the power button to allow wakeup. A full reboot is required.
- **G3**, *Mechanical Off*. The computer's power has been totally removed via a mechanical switch.



Processor Performance States (P-States)

- **P0** max power and frequency
- **P1** less than P0, voltage/frequency scaled
- **P2** less than P1, voltage/frequency scaled
- ...
- **P_n** less than $P(n-1)$, voltage/frequency scaled



Processor “Sleep” States (C-states)

- **C0**: is the operating state.
- **C1** (often known as *Halt*): is a state where the processor is not executing instructions, but can return to an executing state instantaneously. All ACPI-conformant processors must support this power state.
- **C2** (often known as *Stop-Clock*): is a state where the processor maintains all software-visible state, but may take longer to wake up. This processor state is optional.
- **C3** (often known as *Sleep*) is a state where the processor does not need to keep its cache, but maintains other state. This processor state is optional.



Turning Processors Off

The Cost of Wakeup

- Energy expended on wakeup, E_{wake}
- To sleep or not to sleep?



Turning Processors Off

The Cost of Wakeup

- Energy expended on wakeup, E_{wake}
- To sleep or not to sleep?

- Not to sleep (for time t):

$$E_{no-sleep} = (k_v V^2 f + R_0) t$$

- To sleep (for time t) then wake up:

$$E_{sleep} = P_{sleep} t + E_{wake}$$



Turning Processors Off

The Cost of Wakeup

- Energy expended on wakeup, E_{wake}
- To sleep or not to sleep?

- Not to sleep (for time t):

$$E_{no-sleep} = (k_v V^2 f + R_0) t$$

- To sleep (for time t) then wake up:

$$E_{sleep} = P_{sleep} t + E_{wake}$$

- To save energy by sleeping: $E_{sleep} < E_{no-sleep}$

$$t > \frac{E_{wake}}{k_v V^2 f + R_0 - P_{sleep}}$$



Turning Processors Off

The Cost of Wakeup

- Energy expended on wakeup, E_{wake}
- To sleep or not to sleep?

- Not to sleep (for time t):

$$E_{no-sleep} = (k_v V^2 f + R_0) t$$

- To sleep (for time t) then wake up:

$$E_{sleep} = P_{sleep} t + E_{wake}$$

- To save energy by sleeping: $E_{sleep} < E_{no-sleep}$

$$t > \frac{E_{wake}}{k_v V^2 f + R_0 - P_{sleep}}$$

**Minimum sleep
interval**





Dynamic Power Management

- DPM refers to turning devices off (or putting them in deep sleep modes)
- Device wakeup has a cost that imposes a minimum sleep interval (a breakeven time)
- DPM must maximize power savings due to sleep while maintaining schedulability

DPM and the Problem with Work-conserving Scheduling

- Example:

Task 1 ($C=2$, $P=12$)



Task 2 ($C=1$, $P=16$)



DPM and the Problem with Work-conserving Scheduling

■ Example:

Task 1 ($C=2$, $P=12$)



Task 2 ($C=1$, $P=16$)



Minimum sleep period

DPM and the Problem with Work-conserving Scheduling

■ Example:

Task 1 ($C=2$, $P=12$)



Task 2 ($C=1$, $P=16$)



Minimum sleep period


DPM and the Problem with Work-conserving Scheduling

- No opportunity to sleep ☹️

Task 1 (C=2, P=12)



Task 2 (C=1, P=16)



Minimum sleep period

DPM and the Problem with Work-conserving Scheduling

- Must batch! 😊

Task 1 ($C=2, P=12$)




Task 2 ($C=1, P=16$)




Minimum sleep period

A Schedulability Question:

How to Analyze Schedules with Sleep Periods?



Task 1 ($C=2$, $P=12$)



Task 2 ($C=1$, $P=16$)



Minimum sleep period

A Schedulability Question:

How to Analyze Schedules with Sleep Periods?

- Option 1: Treat sleep periods like a sporadic task. Use the Liu and Layland utilization bound for schedulability. Problems?

Task 1 ($C=2, P=12$)



Task 2 ($C=1, P=16$)



Task 3 ($C=11, P=16$)



A Schedulability Question:

How to Analyze Schedules with Sleep Periods?

- Option 1: Treat sleep periods like a periodic task. Use the Liu and Layland utilization bound for schedulability. Problems?
 - Does not work because the “sleep task” cannot be preempted, whereas the rest of the tasks are preemptible. The utilization bound works only for fully preemptive scheduling.

Task 1 ($C=2$, $P=12$)



Task 2 ($C=1$, $P=16$)



Task 3 ($C=11$, $P=16$)



A Schedulability Question:

How to Analyze Schedules with Sleep Periods?

- Option 2: Treat sleep periods like the *highest-priority* periodic task. Use the Liu and Layland utilization bound for schedulability. Problems?

Task 1 ($C=2, P=12$)



Task 2 ($C=1, P=16$)



Task 3 ($C=11, P=16$)



A Schedulability Question:

How to Analyze Schedules with Sleep Periods?

- Option 2: Treat sleep periods like the *highest-priority* periodic task. Use the Liu and Layland utilization bound for schedulability. Problems?
 - Does not work because the “sleep task” may need to have a larger period than the actual top-priority task, which contradicts rate-monotonic scheduling. The bound does not work.

Task 1 ($C=2, P=12$)



Task 2 ($C=1, P=16$)



Task 3 ($C=11, P=16$)



A Schedulability Question:

How to Analyze Schedules with Sleep Periods?

- Option 3: Treat sleep periods like the *highest-priority* periodic task. Use *exact response time analysis* for schedulability. Problems?

Task 3 (C=11, P=16)



Task 1 (C=2, P=12)



Task 2 (C=1, P=16)



Device Forbidden Regions

- Option 3: Treat sleep periods like the *highest-priority* periodic task. Use *exact response time analysis* for schedulability. Problems?


- A Valid solution, but pessimistic.

(Called: Device Forbidden Regions. Published in RTAS 2008.)

Task 3 (C=11, P=16)




Task 1 (C=2, P=12)



Task 2 (C=1, P=16)





Intel CPU Clock Speed



- Moore's Law (1980-2005)
- Question: Why did the speed curve saturate (around 2005)?

Computational Power (per Die)




- Note the exponential rise in power consumption
- Question: how come it does not saturate?

Microprocessor Transistor Counts 1971-2011 & Moore's Law

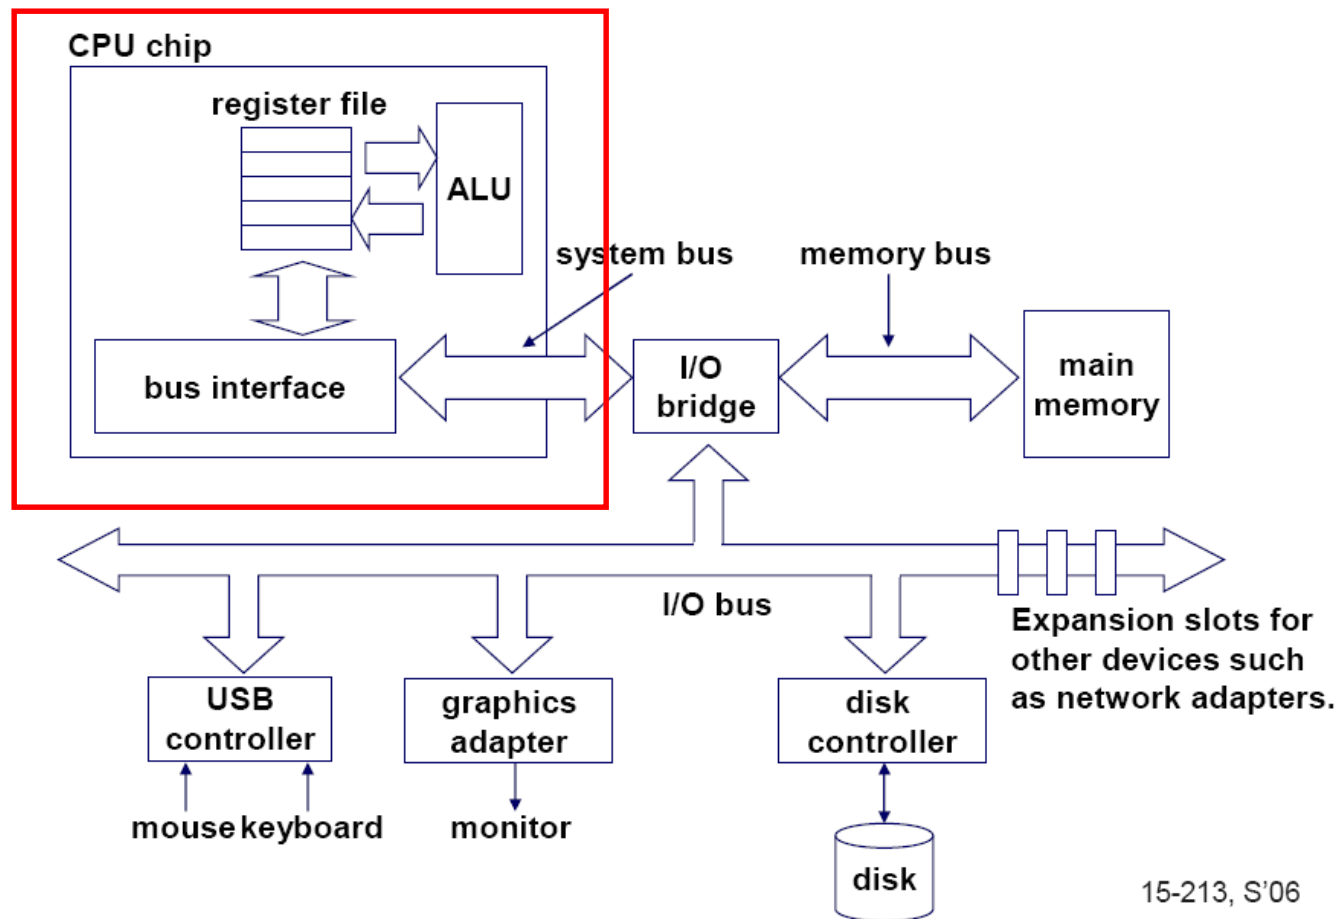
Moore's Law

Transistor count
doubles every two
years

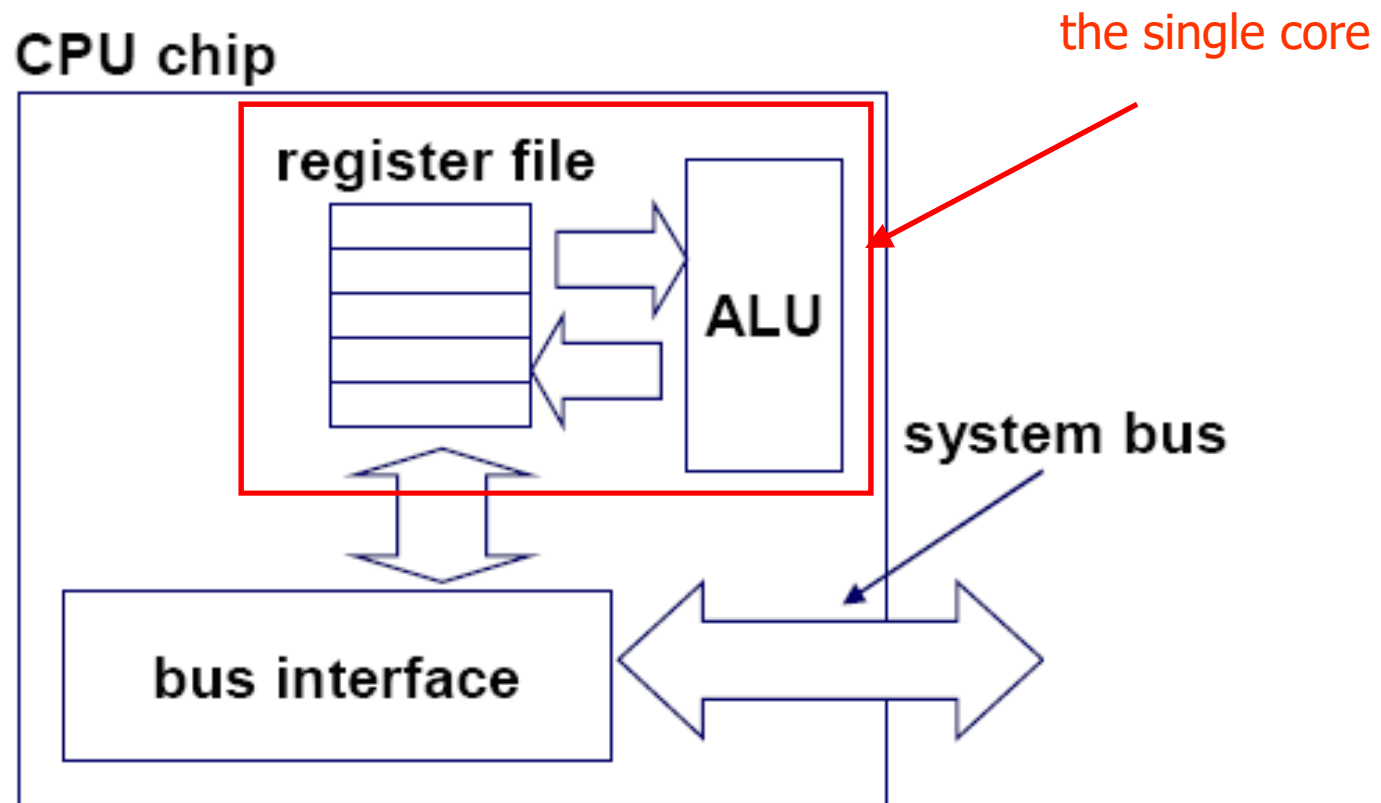


Named after Intel co-founder Gordon E. Moore, who described the trend in his 1965 paper

Single-core computer

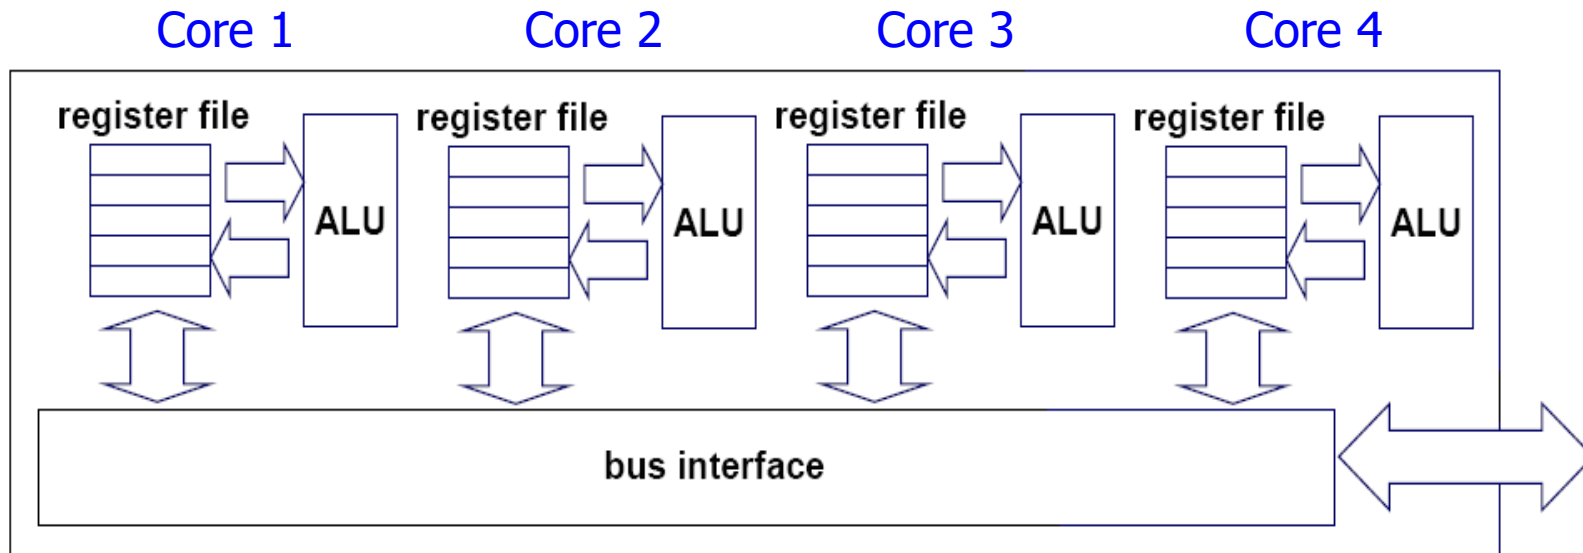


Single-core CPU chip

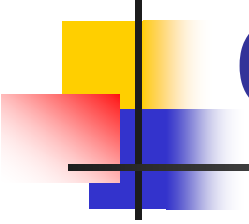


Multi-core Architectures

- Replicate multiple processor cores on a single die.



Multi-core CPU chip



Interaction with the Operating System

- OS perceives each core as a separate processor
- OS scheduler maps threads/processes to different cores
- Major OSes support multi-core today: Windows, Linux, Mac OS X, ...



DVS on Multiprocessors

- Consider a set of tasks, where task i has period P_i and total number of cycles C_i
 - Sort tasks from largest to smallest utilization C_i / P_i
 - Assign tasks one at a time (largest-first) to the least utilized processor
 - Apply one of the previous algorithms on each processor separately



Question

- From the perspective of minimizing energy, is it always a good idea to use up all processors?



How Many Processors to Use?

- Consider using one processor at frequency f versus two at frequency $f/2$
- Case 1: Total power for one processor
 - $k_f f^3 + R_0$
- Case 2: Total power for two processors
 - $2 \{k_f (f/2)^3 + R_0\} = k_f f^3 / 4 + 2 R_0$



How Many Processors to Use?

- Consider using one processor at frequency f versus two at frequency $f/2$
- Case 1: Total power for one processor
 - $k_f f^3 + R_0$
- Case 2: Total power for two processors
 - $2 \{k_f (f/2)^3 + R_0\} = k_f f^3 / 4 + 2 R_0$
- The general case: n processors
 - $n \{k_f (f/n)^3 + R_0\} = k_f f^3 / n^2 + n R_0$



How Many Processors to Use?

- The general case: n processors
 - $Power = n \{k_f (f/n)^3 + R_0\} = k_f f^3 / n^2 + n R_0$
 - $dPower/dn = -2 k_f f^3 / n^3 + R_0 = 0$

$$n = \sqrt[3]{\frac{2k_f f^3}{R_0}}$$



How Many Processors to Use?

- The general case: n processors

- $Power = n \{k_f (f/n)^3 + R_0\} = k_f f^3 / n^2 + n R_0$

- $dPower/dn = -2 k_f f^3 / n^3 + R_0 = 0$

$$n = \sqrt[3]{\frac{2k_f f^3}{R_0}}$$

- What if n is not an integer?