




# Midterm Review #2

---

# Exact Schedulability Test: Example

$$I = \sum_j \left\lceil \frac{R_i}{P_j} \right\rceil C_j$$

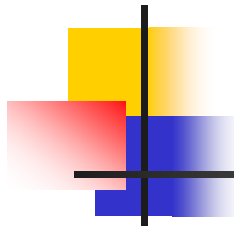
$$R_i = I + C_i$$



Consider a system of two tasks:

Task 1:  $P_1=1.7$ ,  $D_1=0.5$ ,  $C_1=0.5$


Task 2:  $P_2=8$ ,  $D_2=3.2$ ,  $C_2=2$



# Example

$$I = \sum_j \left\lceil \frac{R_i}{P_j} \right\rceil C_j$$

$$R_i = I + C_i$$



Consider a system of two tasks:

Task 1:  $P_1=1.7$ ,  $D_1=0.5$ ,  $C_1=0.5$

Task 2:  $P_2=8$ ,  $D_2=3.2$ ,  $C_2=2$

$$I^{(0)} = C_1 = 0.5$$

$$R_2^{(0)} = I^{(0)} + C_2 = 2.5$$

$$I^{(1)} = \left\lceil \frac{R_2^{(0)}}{P_1} \right\rceil C_1 = \left\lceil \frac{2.5}{1.7} \right\rceil 0.5 = 1$$

$$R_2^{(1)} = I^{(1)} + C_2 = 3$$

$$I^{(2)} = \left\lceil \frac{R_2^{(1)}}{P_1} \right\rceil C_1 = \left\lceil \frac{3}{1.7} \right\rceil 0.5 = 1$$

$$R_2^{(2)} = I^{(2)} + C_2 = 3$$

$3 < 3.2 \rightarrow \text{Ok!}$



# Mixed Periodic and Aperiodic Task Systems

---

- Question: how to execute aperiodic tasks without violating schedulability guarantees given to periodic tasks?




# Polling Server

---

- Runs as a periodic task (priority set according to RM)
- Aperiodic arrivals are queued until the server task is invoked
- When the server is invoked it serves the queue until it is empty or until the budget expires then suspends itself
  - If the queue is empty when the server is invoked it suspends itself immediately.
- Server is treated as a regular periodic task in schedulability analysis


# Example of a Polling Server

- Polling server:
  - Period  $P_s = 5$
  - Budget  $B_s = 2$
- Periodic task
  - $P = 4$
  - $C = 1.5$
- All aperiodic arrivals have  $C=1$



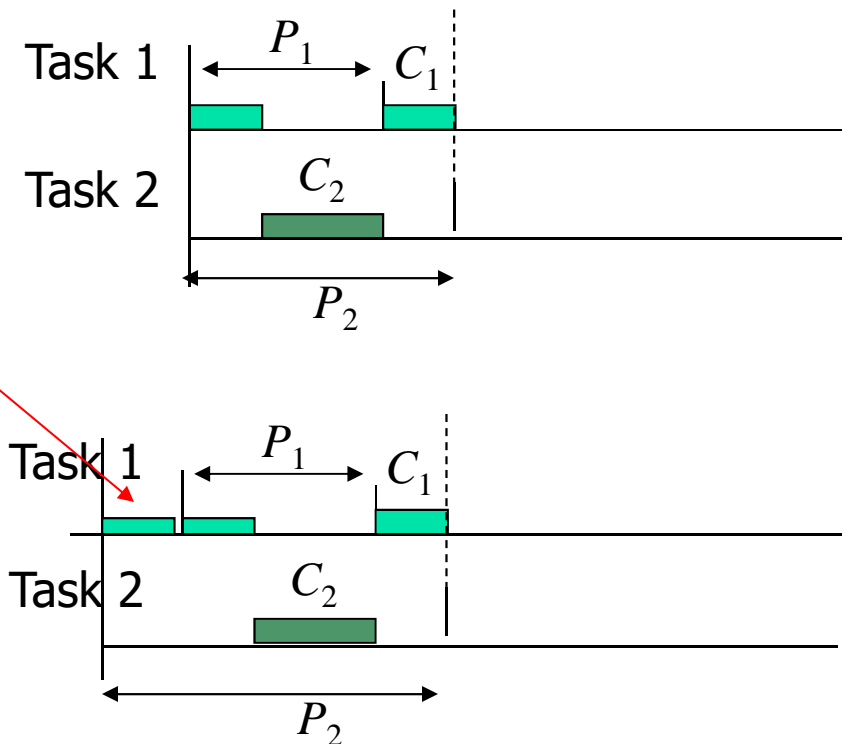
# Deferrable Server

- Keeps the balance of the budget until the end of the period
- Example (continued)



# Worst-Case Scenario

**Deferred  
Previous  
Invocation**



$$U_p \leq \ln \left( \frac{U_s + 2}{2U_s + 1} \right)$$

Exercise: Derive the utilization bound for a deferrable server plus one periodic task





# Priority Exchange Server

---

- Like the deferrable server, it keeps the budget until the end of server period
- Unlike the deferrable server the priority slips over time: When not used the priority is exchanged for that of the executing periodic task


# Priority Exchange Server

## Example

Aperiodic tasks

Priority Exchange  
Server

Periodic  
Tasks



$$U_p \leq \ln\left(\frac{2}{U_s + 1}\right)$$



# Sporadic Server

---

- Server is said to be *active* if it is in the *running* or *ready* queue, otherwise it is *idle*.
- When an aperiodic task comes and the budget is not zero, the server becomes active
- Every time the server becomes *active*, say at  $t_A$ , it sets replenishment time one period into the future,  $t_A + P_s$  (but does not decide on replenishment amount).
- When the server becomes idle, say at  $t_I$ , set replenishment amount to capacity consumed in  $[t_A, t_I]$

$$U_p \leq \ln \left( \frac{2}{U_s + 1} \right)$$



# Slack Stealing Server

---

- Compute a slack function  $A(t_s, t_f)$  that says how much total slack is available
- Admit aperiodic tasks while slack is not exceeded



# Power of Computation

---

- Terminology

- $R$  : Power spent on computation
- $V$  : Processor voltage
- $f$  : Processor clock frequency
- $R_0$  : Leakage power

- Power spent on computation is:

- $R = k_v V^2 f + R_0$   
where  $k_v$  is a constant



# Energy of Computation

---

- Power spent on computation is:
  - $R = k_v V^2 f + R_0$
- Consider a task of length  $C$  clock cycles and a processor operating at frequency  $f$
- The execution time is  $t = C/f$
- Energy spent is:
  - $E = R t = (k_v V^2 f + R_0)(C/f)$



# Reducing Processor Frequency

## Good or Bad?

---

- Does it make sense to operate the processor at a reduced speed to save energy? Why or why not?

Possible Answer:

$$E = R t = (k_v V^2 f + R_0)(C/f) = k_v V^2 C + R_0 C/f$$

- Conclusion:  $E$  is minimum when  $f$  is maximum.
  - Operate at top speed
- Is this really true? What are the underlying assumptions?

# Dynamic Voltage Scaling (DVS):

## Reducing Voltage and Frequency

---

- Processor voltage can be decreased if clock frequency is decreased
  - Voltage and frequency can be decreased roughly proportionally.
  - In this case (where  $V \sim f$ ):

$$R = k_f f^3 + R_0$$

$$E = (k_f f^3 + R_0)(C/f) = k_f f^2 C + R_0 C/f$$



# Dynamic Voltage Scaling (DVS):

## Reducing Voltage and Frequency

- Processor voltage can be decreased if clock frequency is decreased
  - Voltage and frequency can be decreased roughly proportionally.

$$R = k_f f^3 + R_0$$

$$E = (k_f f^3 + R_0)(C/f) = k_f f^2 C + R_0 C/f$$

- Question: Does reducing frequency (and voltage) increase or decrease total energy spend on a task?

# Dynamic Voltage Scaling (DVS):

## The Critical Frequency

---

- There exists a minimum frequency below which no energy savings are achieved

$$E = k_f f^2 C + R_0 C / f$$

$$dE/df = 2k_f f C - R_0 C / f^2 = 0$$

$$f = \sqrt[3]{\frac{R_0}{2k_f}}$$



# DVS Algorithm 1: Static Voltage Scaling

---

1. Calculate the critical frequency
2. Calculate the minimum frequency at which the task set remains schedulable
  - Example: If EDF is used and the utilization is 60% at the maximum frequency  $f_{max}$ , then the frequency can be decreased to  $0.6 f_{max}$ .
3. Let  $f_{opt}$  be the larger of the above two
4. Operate the system at the smallest frequency at or above  $f_{opt}$ .



# DVS Algorithm 2: Cycle-conserving DVS

---

- What if a task finishes early?
  - Re-compute the utilization based on the reduced execution time.
  - Calculate the minimum frequency at which the task set is schedulable using the new utilization.
  - Update task execution times to the WCET when new invocations are released.

# Practical Consideration:

## Accounting for Off-chip Overhead

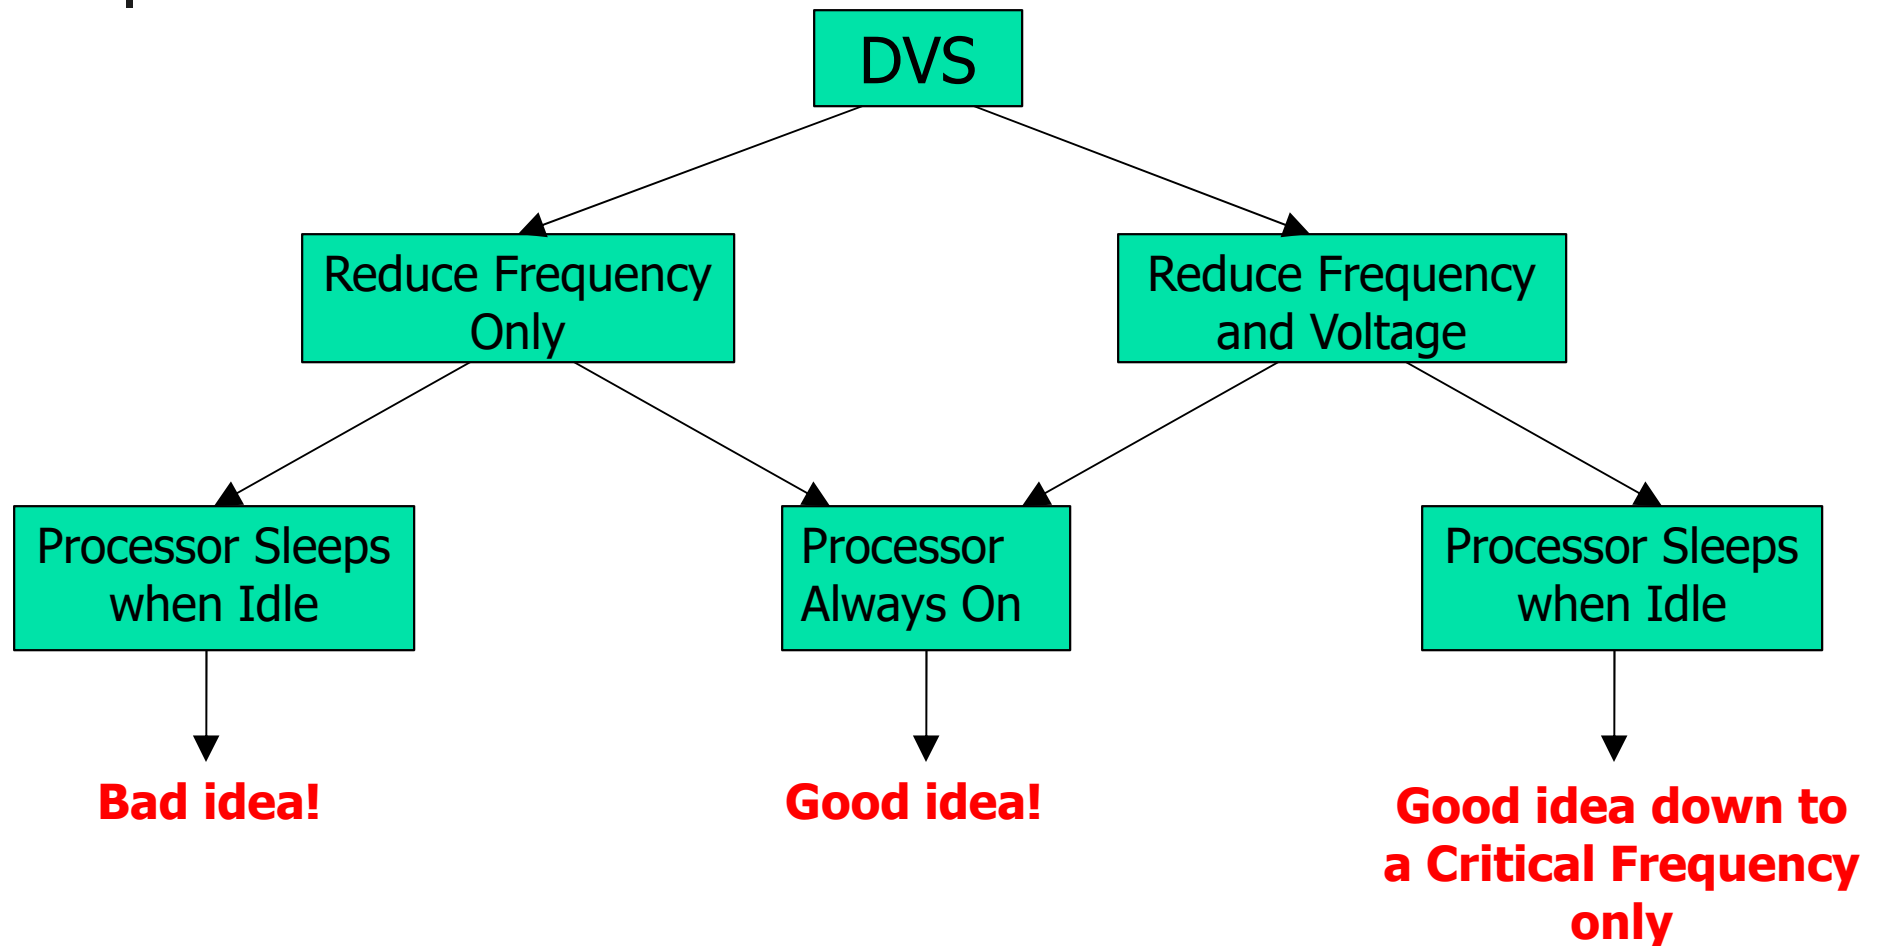
- In the preceding discussion, we assumed that task execution *time* at frequency  $f$  is  $C/f$ , where  $C$  is the total cycles needed
- In reality some cycles are lost waiting for memory access and I/O (Off-chip cycles).
  - Let the number of CPU cycles used be  $C_{cpu}$  and the time spent off-chip be  $C_{off-chip}$
  - Execution time at frequency  $f$  is given by

$$C_{cpu}/f + C_{off-chip}$$



# Recap

---





# Processor Performance States (P-States)

---

- **P0** max power and frequency
- **P1** less than P0, voltage/frequency scaled
- **P2** less than P1, voltage/frequency scaled
- ...
- **P<sub>n</sub>** less than  $P(n-1)$ , voltage/frequency scaled



# Processor “Sleep” States (C-states)

---

- **C0**: is the operating state.
- **C1** (often known as *Halt*): is a state where the processor is not executing instructions, but can return to an executing state instantaneously. All ACPI-conformant processors must support this power state.
- **C2** (often known as *Stop-Clock*): is a state where the processor maintains all software-visible state, but may take longer to wake up. This processor state is optional.
- **C3** (often known as *Sleep*) is a state where the processor does not need to keep its cache, but maintains other state. This processor state is optional.





# Turning Processors Off

## The Cost of Wakeup

---

- Energy expended on wakeup,  $E_{wake}$
- To sleep or not to sleep?

- Not to sleep (for time  $t$ ):

$$E_{no-sleep} = (k_v V^2 f + R_0) t$$

- To sleep (for time  $t$ ) then wake up:

$$E_{sleep} = P_{sleep} t + E_{wake}$$

- To save energy by sleeping:  $E_{sleep} < E_{no-sleep}$

$$t > \frac{E_{wake}}{k_v V^2 f + R_0 - P_{sleep}}$$


# DPM and the Problem with Work-conserving Scheduling

- No opportunity to sleep ☹️

Task 1 (C=2, P=12)



Task 2 (C=1, P=16)




Minimum sleep period

# DPM and the Problem with Work-conserving Scheduling

- Must batch! 😊

Task 1 ( $C=2, P=12$ )



Task 2 ( $C=1, P=16$ )



Minimum sleep period

# Device Forbidden Regions

- Treat sleep periods like the *highest-priority* sporadic task. Use *response time analysis* for schedulability. Problems?
  - A Valid solution, but pessimistic.  
(Called: Device Forbidden Regions. Published in RTAS 2008.)

Task 3 (C=11, P=16)



Task 1 (C=2, P=12)



Task 2 (C=1, P=16)





# How Many Processors to Use?

---

- Consider using one processor at frequency  $f$  versus two at frequency  $f/2$
- Case 1: Total power for one processor
  - $k_f f^3 + R_0$
- Case 2: Total power for two processors
  - $2 \{k_f (f/2)^3 + R_0\} = k_f f^3 / 4 + 2 R_0$
- The general case:  $n$  processors
  - $n \{k_f (f/n)^3 + R_0\} = k_f f^3 / n^2 + n R_0$



# How Many Processors to Use?

---

- The general case:  $n$  processors
  - $Power = n \{k_f (f/n)^3 + R_0\} = k_f f^3 / n^2 + n R_0$
  - $dPower/dn = -2 k_f f^3 / n^3 + R_0 = 0$

$$n = \sqrt[3]{\frac{2k_f f^3}{R_0}}$$



# Relation of Temperature and Energy

---

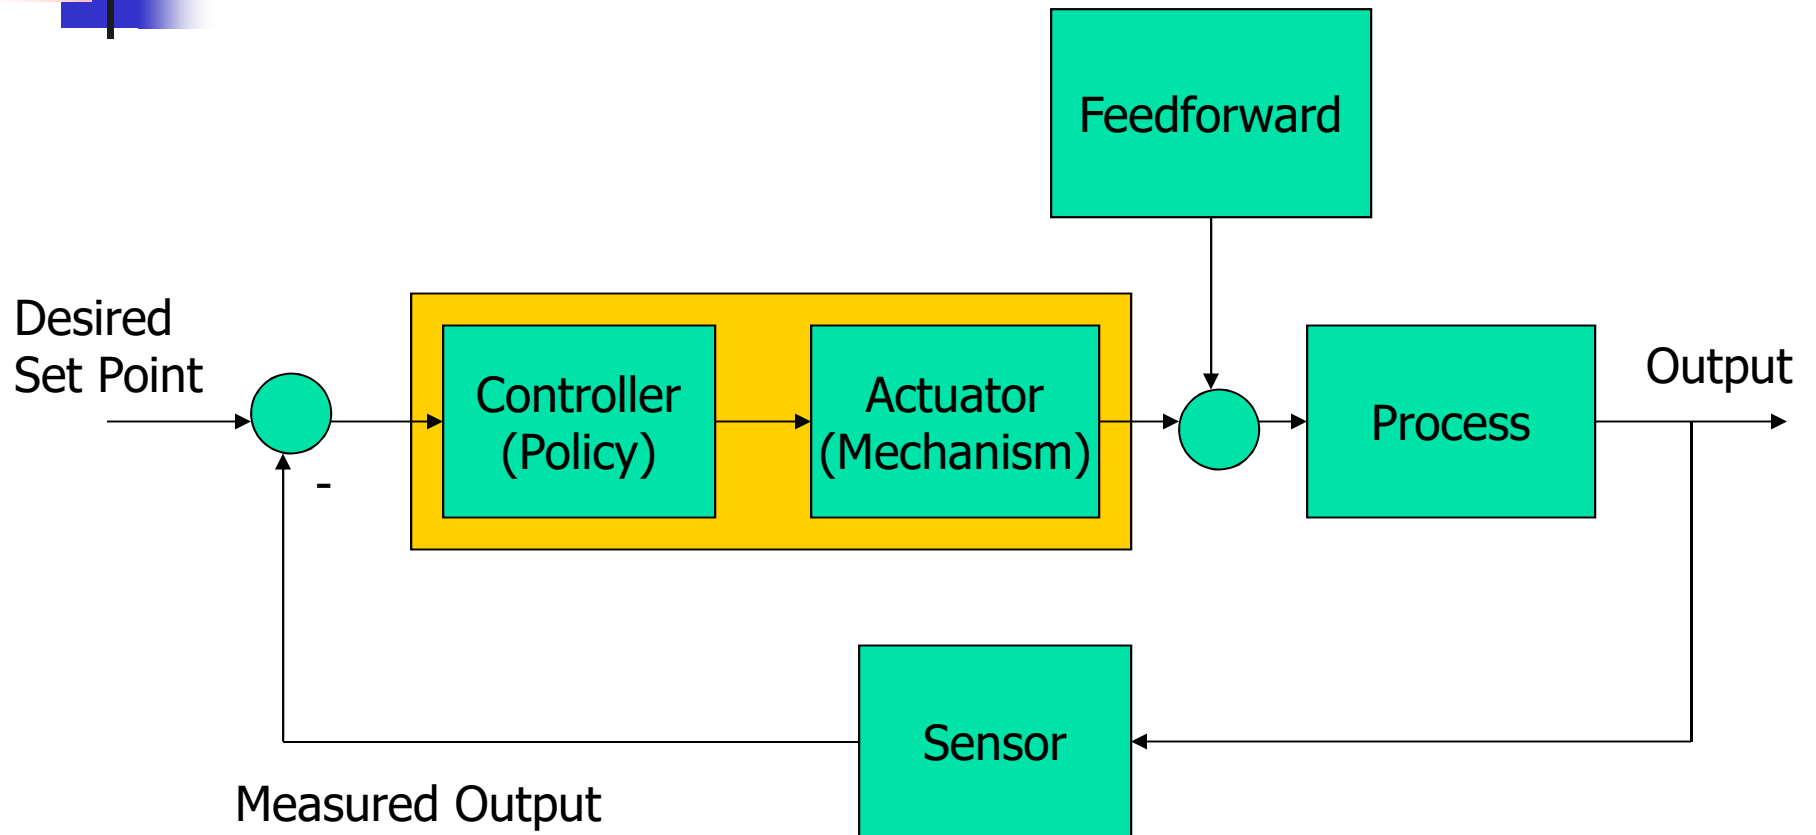
- The rate of change of temperature is proportional to the difference between input power and output power (via cooling)

$$\frac{dT}{dt} = P_{in} - P_{out}$$

$$P_{in} = f(DVS, sleep)$$

$$P_{out} = g(T)$$

# Classical Feedback Control Loops






# Stability – Recap

Phase equation:  $\sum_i p_i(f) = \pi$   $\longrightarrow$   $f$  is obtained

Gain equation:  $\prod_i g_i(f)$  must be less than 1 for stability





# Summary of Basic Elements

Input =  $\sin(\omega t)$

Element	Gain	Phase
Integrator	$1/\omega$	$-\pi/2$
Differentiator	$\omega$	$\pi/2$
Pure delay element (Delay = $D$ )	1	$-\omega D$
First order lag (time constant = $\tau$ )	$K / \sqrt{1 + (\tau\omega)^2}$	$-\tan^{-1}(\omega\tau)$
Pure gain (Gain = $K$ )	$K$	0

Note:

$$\omega = 2\pi f_{osc}$$


Where  $f_{osc}$  is  
the loop  
frequency of  
oscillation

# Steady State Error

At steady state the system “catches up” – phase shift is zero.

$$T_r - e \cdot g_1(f) \cdot g_2(f) \cdot g_3(f) \cdot g_4(f) = e$$

$$e = \frac{T_r}{1 + \prod_i g_i(f)}$$





# Example


---

- When a processor executes a task, its temperature increases with normalized frequency,  $f$ , with a time constant = 1 minute. An increase of 0.1 in  $f$  causes a temperature increase of 3 degrees. An on-chip sensor reports temperature with a 10 second delay. Design a frequency controller such that processor temperature stays at or around 80 degrees.

# Solution:

## Step 1: Plot the Control Loop

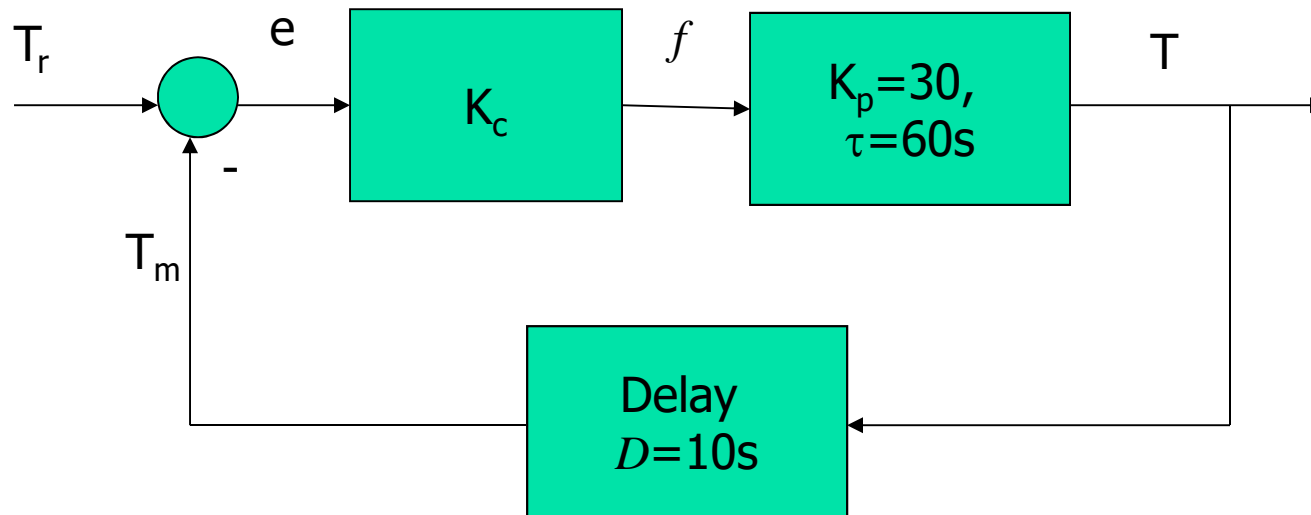
Read the question/narrative carefully and identify all the blocks involved.



# Solution:

## Step 2: Identify the Type and Parameters of all Blocks

Identify all the block types and their parameters.



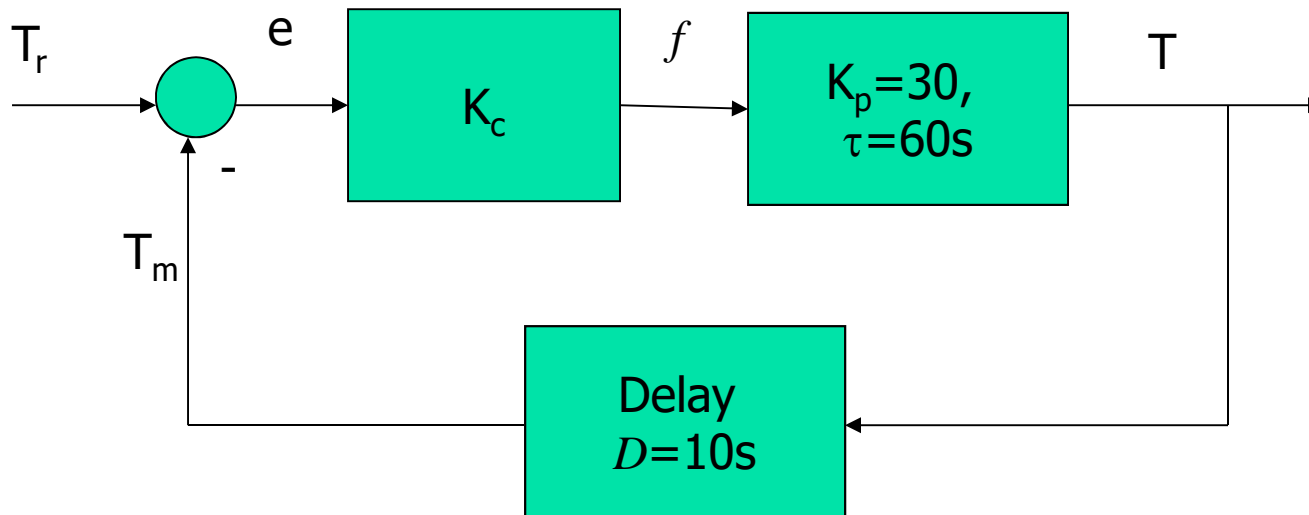
# Solution:

## Step 3: Compute the Natural Frequency of Oscillation from Phase Equation

Substitute in the phase equation and solve for  $f$ .

$$-2\pi f(10) - \tan^{-1}(2\pi f 60) = -\pi$$

Thus,  $f = 0.027$



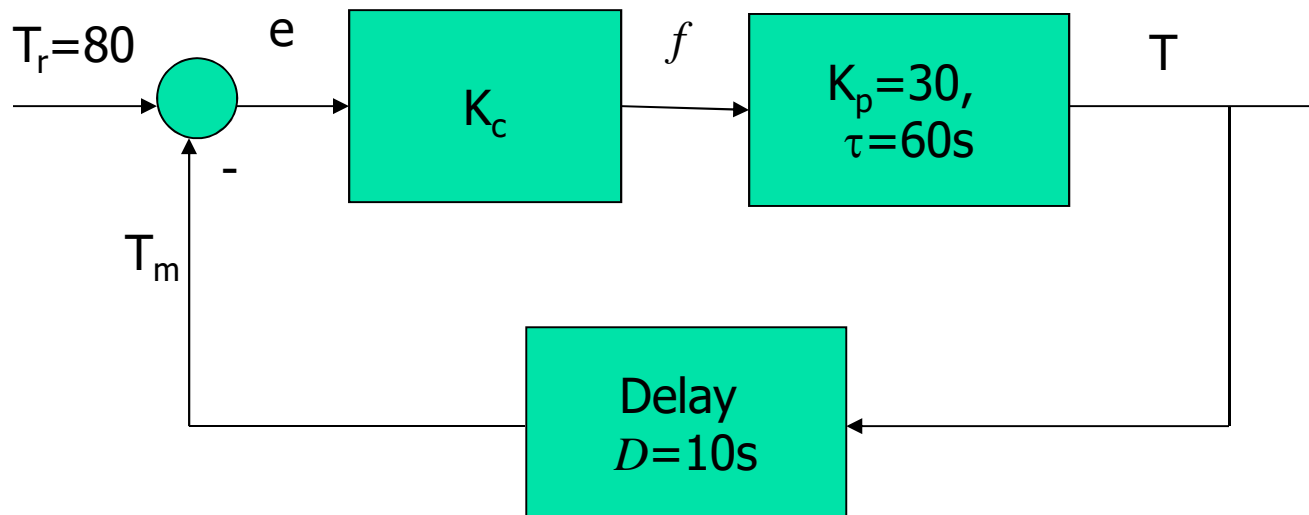
# Solution:

## Step 4: Compute Controller Gain from Gain Equation

Substitute in the gain equation and solve for  $K_c$

$$K_c \cdot 30 / \text{Sqrt} (1 + (2\pi f \cdot 60)^2) = 0.5 \quad \text{where } f = 0.027$$

$$\text{Thus, } K_c = 0.17$$



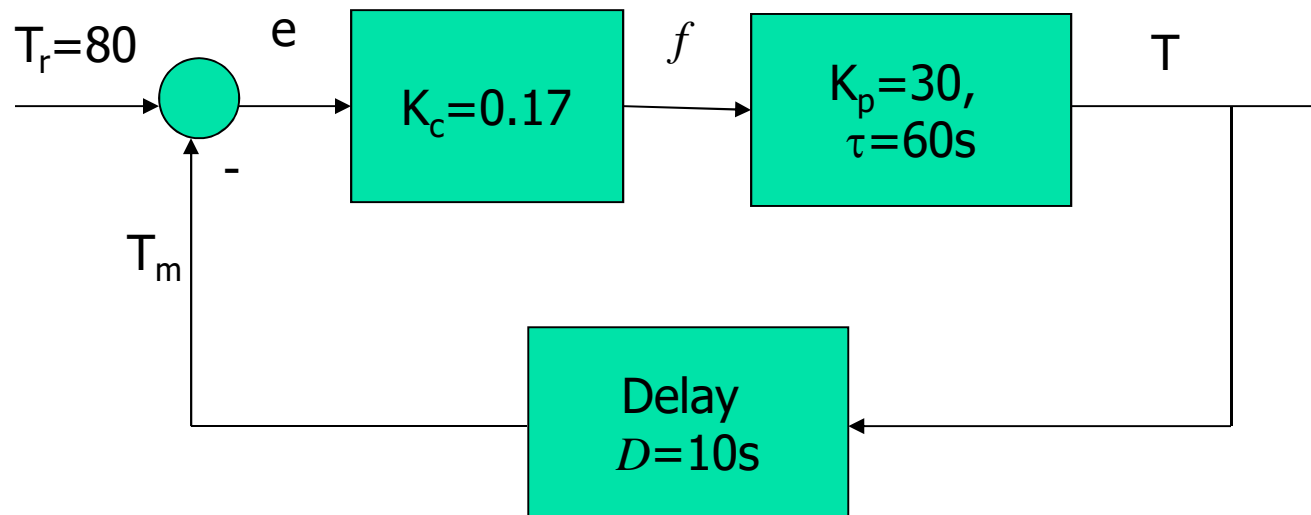
In other words, the control policy should set the processor frequency  $f$  to  $0.17 e$ , where  $e$  is the difference between 80 and the actual processor temperature.



# Solution:

## Note: Steady State Error Calculation

The steady state error is  $80/(1 + 30 \cdot 0.17) = 13$  degrees.





# Example

---

- When a processor executes a task, its temperature increases with normalized frequency,  $f$ , with a time constant = 1 minute. An increase of 0.1 in  $f$  causes a temperature increase of 3 degrees. An on-chip sensor reports temperature with a 10 second delay. Design a frequency controller such that processor temperature stays at or around 80 degrees.
- How would your design change if the energy-optimal frequency of the processor was 0.5?



# Example

---

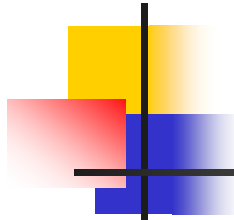
- When a processor executes a task, its temperature increases with normalized frequency,  $f$ , with a time constant = 1 minute. An increase of 0.1 in  $f$  causes a temperature increase of 3 degrees. An on-chip sensor reports temperature with a 10 second delay. Design a frequency controller such that processor temperature stays at or around 80 degrees.
- How would your design change if the energy-optimal frequency of the processor was 0.5?
- How would your design change if in addition the processor executed real-time tasks using RM?



# Example

---

- When a processor executes a task, its temperature increases with normalized frequency,  $f$ , with a time constant = 1 minute. An increase of 0.1 in  $f$  causes a temperature increase of 3 degrees. An on-chip sensor reports temperature with a 10 second delay. Design a frequency controller such that processor temperature stays at or around 80 degrees.
- How would your design change if the energy-optimal frequency of the processor was 0.5?
- How would your design change if in addition the processor executed real-time tasks using RM?
- Does your control loop ensure zero steady state error? If not, redesign.



# Example

---

- When a processor executes a task, its temperature increases with normalized frequency,  $f$ , with a time constant = 1 minute. An increase of 0.1 in  $f$  causes a temperature increase of 3 degrees. An on-chip sensor reports temperature with a 10 second delay. Design a frequency controller such that processor temperature stays at or around 80 degrees.
- How would your design change if the energy-optimal frequency of the processor was 0.5?
- How would your design change if in addition the processor executed real-time tasks using RM?
- Does your control loop ensure zero steady state error? If not, redesign.
- What is the effect of sensor delay on frequency of oscillation?